

CS411 Project Track 2 Final Report

Team No. 8

Team Name: Quartet

Team Member:

Zhoutao Pei
zpei3

Yuetong Liu
liu298

Shuomeng Guang
sguang2

Zhuang Zuo
zhuangz3

1. The Architecture of SimpleDB System

The original and extended SimpleDB system are shown in Fig 1. The gray blocks belongs to the original system, the blue ones are revised classes, and the yellow ones are extended new functions.

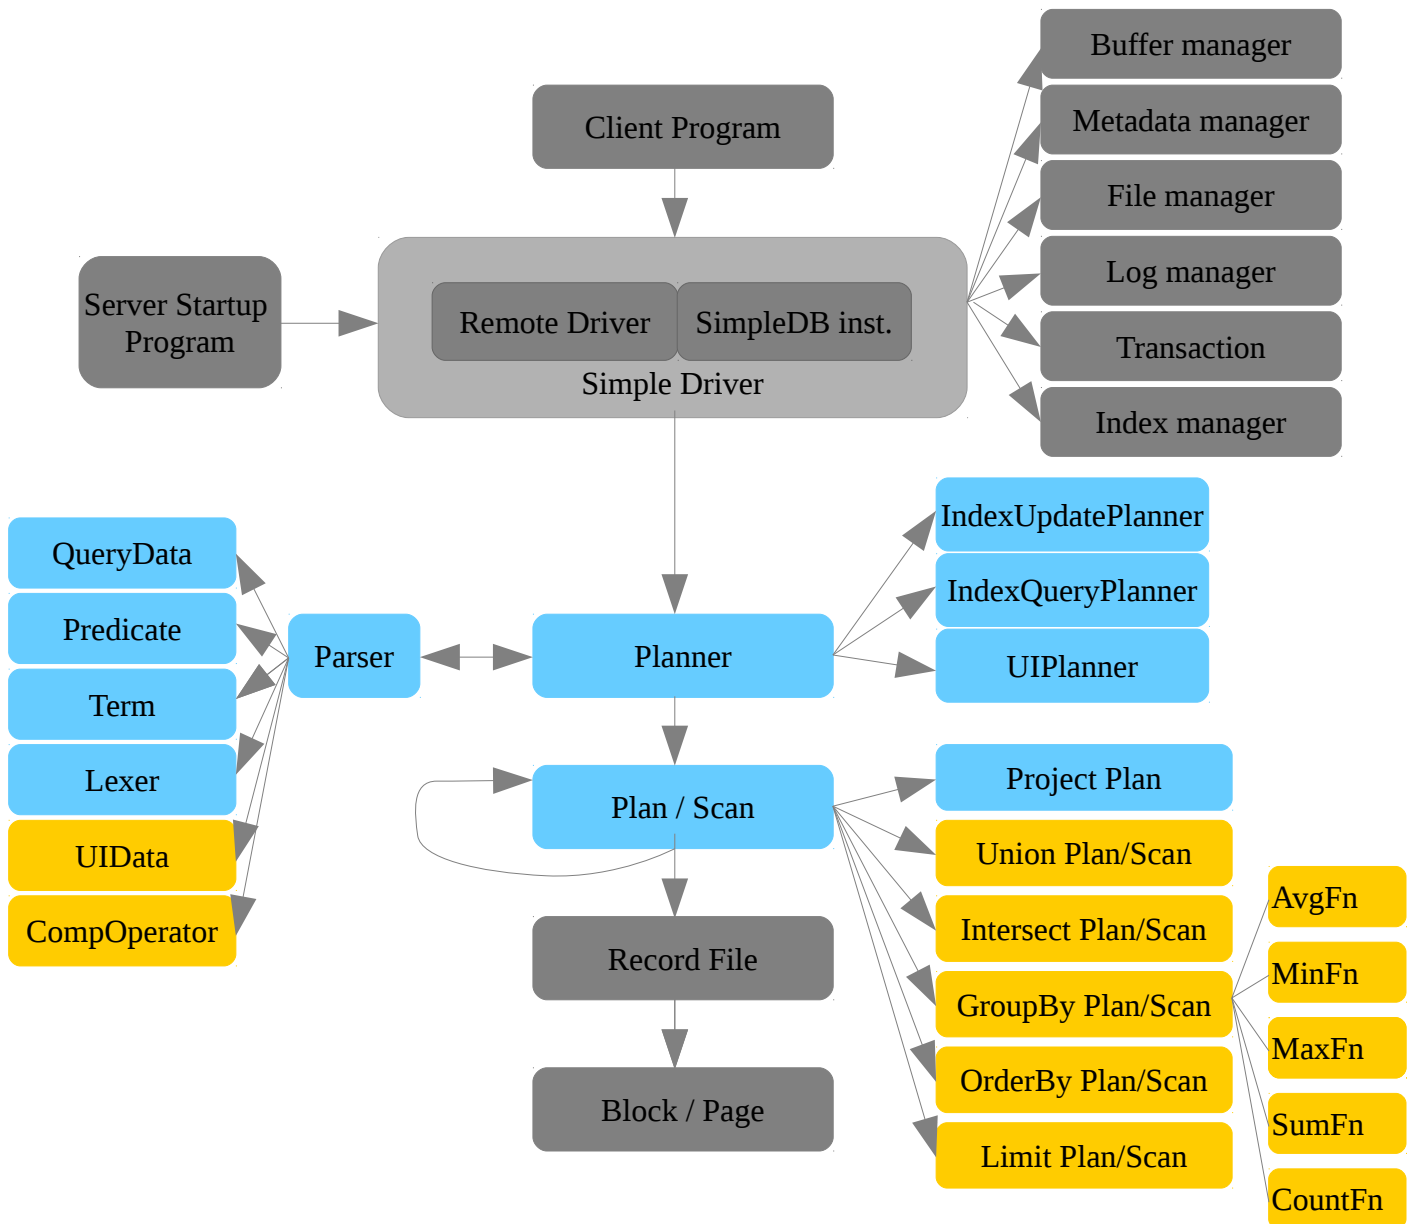


Fig 1. The Architecture of SimpleDB System

2. Function Implementation

2.1 Preparation

To test the new functions of our SimpleDB system, you need to compile it and create a database using the following code:

```
cd SimpleDB_2.10
javac -cp . simpledb/**/*.java simpledb/**/*.java      #compile SimpleDB
javac -cp . studentClient/simpledb/*.java                 #compile client code
java simpledb.server.Startup studentdb
jar cf simpledb.jar simpledb/**/*.class simpledb/**/*.class
cp simpledb.jar studentClient/simpledb/
cd studentClient/simpledb
java -cp simpledb.jar:. CreateStudentDB                   #create a database and add some tables
java -cp simpledb.jar:. AddTable                          # add some larger tables (to test the efficiency of index)
# This might take about 1 minute
# To run AddTable, you need 3 txt files, "direct.txt", "director.txt" and "movie.txt", which
# are in folder studentClient/simpledb/
java -cp simpledb.jar:. SQLInterpreter                    # now you can run queries in console
```

2.2 Comparison operator

The original SimpleDB only support predicate in the form “A = B”. We added more comparison operators, including “<”, “>”, “<=”, “>=” and “<>”. These operators can be used to compare both integers and strings.

Test case:

```
select majorid, gradyear, sname, sid from student where gradyear >= 2005
```

Query result:

majorid	gradyear	sname	sid
10	2005	max	3
20	2005	sue	4

Test case:

```
select majorid, gradyear, sname, sid from student where sname < 'max'
```

Query result:

majorid	gradyear	sname	sid
10	2004	joe	1
20	2004	amy	2
30	2003	bob	5
20	2001	kim	6
30	2004	art	7
10	2004	lee	9

Files modified or added:

Parser.java
revised: method term()
Term.java

revised: constructor, method isSatisfied()

added: member variable Oper

CompOperator.java : new class (under query package)

2.3 Select all

We added support for “*” to allow selecting all fields from the tables. To implement this feature, we revised Parser, so that it will return an empty field list if it sees “*”, and ProjectPlan will extract all the fields from schema if the field list is empty.

Test case:

*select * from student, dept where majorid = did and sid < 3*

Query result:

<i>majorid</i>	<i>gradyear</i>	<i>sname</i>	<i>dname</i>	<i>did</i>	<i>sid</i>
10	2004	joe	compsci	10	1
20	2004	amy	math	20	2

Files modified or added:

ProjectPlan.java

revised: constructor

Parser.java

revised: method query()

QueryData.java

revised: method toString()

StringConstant.java

revised: method toString()

2.4 Union and intersect

Since union and intersect operation will combine two query results, we wrote a UnionPlanner instead of modify any existing planner. We also wrote new plan and scan. The basic idea for implementing union is to maintain a set in UnionScan to remember the value it has seen, and for intersect, this set will read in records from one table first and then look for same records in the other table. Which table to read in first depends on which table has less records. The function we implemented are set union and set intersect, so you won't see duplicates in the outputs.

Test case:

*select * from student where sid < 2 union select * from student where sid > 8*

Query result:

<i>majorid</i>	<i>gradyear</i>	<i>sname</i>	<i>sid</i>
10	2004	joe	1
10	2004	lee	9

Test case:

*select * from student where sid > 2 intersect select * from student where sid < 8*

Query result:

<i>majorid</i>	<i>gradyear</i>	<i>sname</i>	<i>sid</i>
10	2005	max	3
20	2005	sue	4
30	2003	bob	5
20	2001	kim	6
30	2004	art	7

Files modified or added:

Planner.java
 revised: constructor, method createQueryPlan()
 added: member variable uiplanner,
 SimpleDB.java
 revised: method planner()
 Lexer.java
 revised: keywords
 Parser.java
 added method unionIntersect()
 UIData.java: new class, under parse package
 UIPlanner.java: new class, under planner package
 UnionPlan.java: new class, under query package
 UnionScan.java: new class, under query package
 IntersectPlan.java: new class, under query package
 IntersectScan.java: new class, under query package

2.5 Complex predicate

The original SimpleDB only supports multiple term connected by keyword “and”. For each record, it would check each of the term, if any one of them is not satisfied, then the record is skipped.

In our extended SimpleDB system, the predicate can be more complex. We added support for “or” and also parenthesis to separate terms in query predicate.

To implement complex predicate, we added three arraylists in predicate to save the information for left, right parenthesis and logical operators. Basically, we used two stacks to evaluate if a predicate is true or not, Fig 2 shows a simple example:

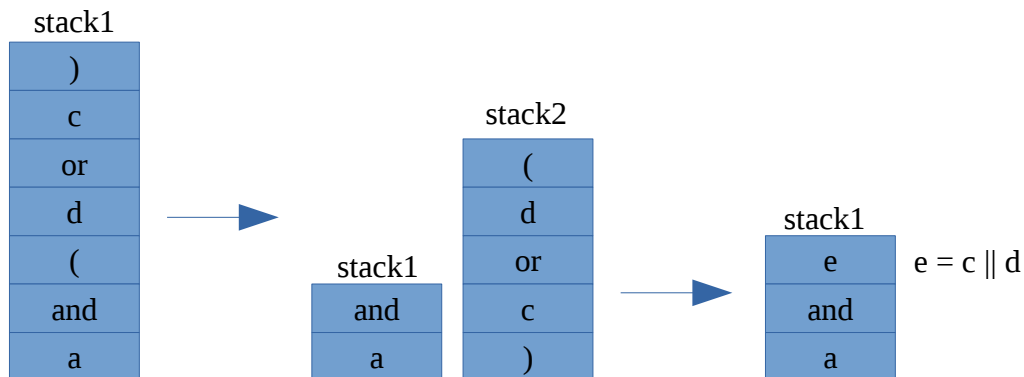


Fig 2. Use 2 stacks to evaluate a complex predicate

We also implemented “in” operator by converting “A in (a, b, c)” into a series of “or” terms “A = a or A = b or A = c”.

Test case:

*select * from student, enroll where ((sid=studentid) and ((gradyear in (2004,2005)) or (grade='A')))*

Query result:

<i>studentid</i>	<i>eid</i>	<i>majorid</i>	<i>gradyear</i>	<i>sname</i>	<i>grade</i>	<i>sectionid</i>	<i>sid</i>
1	14	10	2004	joe	A	13	1
1	24	10	2004	joe	C	43	1
2	34	20	2004	amy	B+	43	2
4	44	20	2005	sue	B	33	4
4	54	20	2005	sue	A	53	4
6	64	20	2001	kim	A	53	6

Files modified or added:

Lexer.java

revised: keyword

Parser.java

revised: method term(), predicate()

Predicate.java

revised: method isSatisfied()

added: member variables leftBrac, rightBrac, logicalOper, isComplex, method setComp()
simpleIsSatisfied(), complexIsSatisfied()

2.6 Group by

We utilized the GroupByPlan and GroupByScan under “materialize” folder and add three more aggregation function “avg”, “min” and “sum” apart from “count” and “max”.

Test case:

select count(sid), majorid, gradyear from student group by majorid, gradyear

Query result:

<i>countofsid</i>	<i>majorid</i>	<i>gradyear</i>
2	10	2004
1	10	2005
2	20	2001
1	20	2004
.....		

Files modified or added:

Parser.java

revised: method query(), selectList()

added: method groupList()

QueryData.java

revised: method toString()

added: member variable groupFields, constructor, method aggFns(), groupFields(),
 BasicQueryPlanner.java
 revised: method createPlan()
 Lexer.java
 revised: keyword
 AvgFn.java - new class, under materialize package
 MinFn.java - new class, under materialize package
 SumFn.java - new class, under materialize package

2.7 Order by

We implemented “order by” operator by using materialization. The method open in OrderPlan will preprocess the input records, split them into up to 2 sorted temporary tables. Then they are passed into OrderScan for merging. Each call to method next in OrderScan will simply retrieve the next record from the sorted temporary tables.

In our extended SimpleDB, records can be sorted by ascending or descending order. This is implemented by maintaining a list of booleans in OrderPlan to indicate the sorting order for each column.

Test case:

*select * from student,dept where majorid=did order by gradyear desc, dname*

Query result:

<i>majorid</i>	<i>gradyear</i>	<i>sname</i>	<i>dname</i>	<i>did</i>	<i>sid</i>
10	2005	max	compsci	10	3
20	2005	sue	math	20	4
10	2004	lee	compsci	10	9
10	2004	joe	compsci	10	1
.....					

Files modified or added:

Lexer.java
 revised: keyword
 Parser.java
 revised: method query()
 added: method sortList()
 QueryData.java:
 revised: constructor
 added: member variable sortFields, method sortFields(), desc()
 BasicQueryPlanner
 revised: method createPlan()
 OrderPlan.java - new class, under materialize package
 OrderScan.java - new class, under materialize package
 OrderRecordComparator.java - new class, under materialize package

2.8 Limit

The “limit” function is implemented by adding LimitPlan and LimitScan, which only output a certain number of records.

The most complicated query flow now becomes:

TablePlan → ProductPlan → SelectPlan → GroupByPlan → OrderPlan → ProjectPlan → LimitPlan

Here we give a complex test case:

select sectionid, count(sid) from student, enroll where sid=studentid and (gradyear in (2004,2005) or sid > 5) group by sectionid order by sectionid desc limit 2

Query result:

countofsid	sectionid
2	53
2	43

Files modified or added:

Parser.java

revised: method query()

QueryData.java

added: member variable, constructor, method toString(), add limit()

Lexer.java

revised: keyword

BasicQueryPlanner.java

revised: method createPlan()

LimitPlan.java - new class, under query package

LimitScan.java - new class, under query package

2.9 Index selection, index join and query optimization

“Create index” has been implemented in the original SimpleDB. We revised IndexInfo so that it can use BTree index. To utilize index in query, a new IndexQueryPlanner was created.

To compare the difference between query with index and without index, here we use table director, direct and movie, which are created by program AddTable in section 2.1. The relationship between these tables are shown below:

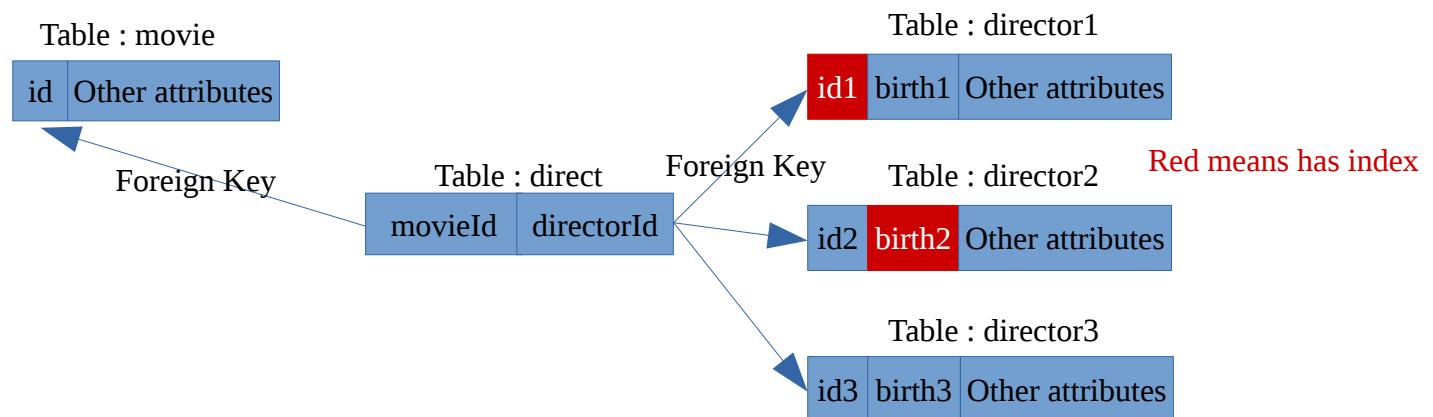


Fig 3. Relationship between tables used for testing index query

Table director1, director2 and director3 has exactly the same content, however, director1 has index for id1, and director2 has index for birth, director3 has no index. We revised SQLInterpreter so that it can show running time for each query.

In IndexQueryPlanner, we didn't follow the sequence of query plans in BasicQueryPlanner. Instead, we did some query optimization. If a term is in the form "Column A = some constant", SelectPlan will be executed directly after TablePlan, instead of ProductPlan and SelectPlan in BasicQueryPlanner. If Column A contains index, IndexSelectPlan will be executed.

Test case:

Query1:

select name2, title from director2, direct, movie where id2 = directorId and id = movieId and birth2 = 1964 limit 5

Query2(*be careful, this might take 1 minute to run*):

select name3, title from director3, direct, movie where id3 = directorId and id = movieId and birth3 = 1964 limit 5

Query result:

<i>name2</i>	<i>title</i>
<i>EytanFox</i>	<i>TheBubble</i>
<i>PeterFlinth</i>	<i>ArnTheKnightTemplar</i>
<i>StephanElliott</i>	<i>EyeoftheBeholder</i>
<i>StephanElliott</i>	<i>EasyVirtue</i>
<i>GuillermodeiToro</i>	<i>PacificRim</i>

Time Elapsed for Query1: 5152832051
Time Elapsed for Query2: 57829743627

The query results are the same, but the running time for query 2 is significantly longer than query1.

In IndexQueryPlanner, if more than two tables are joined, the table with index on joining field will be joined first, and IndexJoinPlan will be performed, instead of ProductPlan and SelectPlan.

Test case:

Query1:

select name1, title from director1, direct, movie where id1 = directorId and id = movieId and year = 1955 limit 5

Query2(**do not** try this, this will take more than 10 mins):

select name3, title from director3, direct, movie where id3 = directorId and id = movieId and year = 1955 limit 5

Query result:

<i>title</i>	<i>name1</i>
<i>Ordet</i>	<i>CarlTheodorDreyer</i>

Time Elapsed for Query1 : 2278075587
Time Elapsed for Query2 : 996650645315

Again, the query results will be the same, but the running time for query 2 is significantly longer than query1.

Files modified or added:

Planner.java

revised: method executeUpdate()

added: method hasIdx(),

Planner.java

revised: method createQueryPlan()

added: method predHasIdx()

IndexQueryPlanner.java - new class, under index.planner package

IndexInfo.java

revised: method open(), blockAccessed()

Predicate.java

added: method size()

3. Lessons Learned

Firstly, none of us have written code in Java before, so this project is quite challenging. At first, we don't even know how to compile the code and run test cases. Through learning Java and getting familiar with the code, we start from the easy part, discuss about the structure, and test with small changes. The learning curve is sharp but worthwhile.

More important, since the database contains multiple units, it is key to figure out the relation between each units. An overview of the architecture is very important.